

В этом ноутбуке рассматриваются разные методы стекинга моделей, полученных в других ноутбуках

```
In [2]: 1 import numpy as np
2 import matplotlib.pyplot as plt
3 import pandas as pd
4 from sklearn.ensemble import RandomForestClassifier
5 from sklearn.model_selection import train_test_split
6 from sklearn.metrics import f1_score
7 from tqdm import tqdm_notebook
8 from sklearn.linear_model import LinearRegression
9 from sklearn.ensemble import RandomForestClassifier
10 from xgboost import XGBRegressor, XGBClassifier
11 from sklearn.model_selection import cross_val_score, KFold
```


```
In [3]: 1 def pred_proc(pred):
2     pred = np.round(np.clip(pred, 0, 10))
3     return pred.astype(int)
4
5 def cv_loop(oof_df, y, model):
6     n_fold = 5
7     folds = KFold(n_splits=n_fold, shuffle=True, random_state=17)
8
9     scores = []
10
11     for training_index, validation_index in tqdm_notebook(folds.split(oof_df, y)):
12         X_train = oof_df.iloc[training_index]
13         y_train = y[training_index]
14         X_valid = oof_df.iloc[validation_index]
15         y_valid = y[validation_index]
16
17         model.fit(X_train, y_train)
18
19         # скор на валидации
20         preds = model.predict(X_valid)
21         preds = np.round(np.clip(preds, 0, 10)).astype(int)
22         score = f1_score(y_valid, preds, average = 'macro')
23         scores.append(score)
24         print(score)
25
26     return scores
```

1. Стекинг LGBM и Catboost

1.1. Усреднение ответов

```
In [4]: 1 lgb_preds = np.load('for_stacking/lgb13_test_preds.npy')
2 cb_preds = np.load('for_stacking/preds_best_catboost.npy')
```

```
In [4]: 1 plt.figure(figsize=(10, 5))
2         plt.subplot(121)
3         plt.hist(lgb_preds)
4         plt.subplot(122)
5         plt.hist(cb_preds)
6         plt.show()
```



```
In [5]: 1 preds = 0.5*lgb_preds + 0.5*cb_preds
2         preds = pred_proc(preds)
3         sample_df = pd.read_csv("data/sample_submission.csv", dtype={'time':str})
4         sample_df['open_channels'] = preds
5         sample_df.to_csv("lgb13_bestcb.csv", index=False, float_format='%.4f')
```

1.2. Метамодел ь - линейная регрессия

```
In [3]: 1 lgb_oof = np.load('for_stacking/lgb13_oof_preds.npy')
2         cb_oof = np.load('for_stacking/oof_best_catboost.npy')
```

```
In [4]: 1 oof_df = pd.DataFrame({'lgb_oof': lgb_oof, 'cb_oof': cb_oof})
2         oof_df.head()
```

Out[4]:

	lgb_oof	cb_oof
0	0.132022	0.030460
1	0.097993	-0.028301
2	0.081194	-0.085979
3	0.138549	0.036849
4	0.060607	0.035619

```
In [9]: 1 train = pd.read_csv('data/train_clean.csv')
2 test = pd.read_csv('data/test_clean.csv')
3 y = train.open_channels.values
4 train = train.drop(columns=['open_channels'])
```

```
In [39]: 1 X_train, X_val, y_train, y_val = train_test_split(oof_df, y,
2                                                    random_state=17,
3                                                    test_size=0.3)
```

```
In [40]: 1 %%time
2 lr = LinearRegression()
3 lr.fit(X_train, y_train)
```

Wall time: 601 ms

Out[40]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)

```
In [41]: 1 preds = lr.predict(X_val)
2 preds = np.round(np.clip(preds, 0, 10)).astype(int)
3 score = f1_score(y_val, preds, average = 'macro')
4 print(f'score on val: {score}')
```

score on val: 0.937807230681399

```
In [42]: 1 scores = cv_loop(oof_df, y, LinearRegression())
2 print(scores)
3 print(np.mean(scores))
```

A Jupyter widget could not be displayed because the widget state could not be found. This could happen if the kernel storing the widget is no longer available, or if the widget state was not saved in the notebook. You may be able to create the widget by running the appropriate cells.

[0.9378667233513109, 0.938132013690785, 0.9376538834963697, 0.9393951023074935, 0.9386472946846514]
0.938339003506122

```
In [25]: 1 test_df = pd.DataFrame({'lgb_oof': lgb_preds, 'cb_oof': cb_preds})
```

```
In [28]: 1 lr.fit(oof_df, y)
2 print(lr.coef_)
3 preds = lr.predict(test_df)
4 preds = np.round(np.clip(preds, 0, 10)).astype(int)
5 sample_df = pd.read_csv("data/sample_submission.csv", dtype={'time':str})
6 sample_df['open_channels'] = preds
7 sample_df.to_csv('lgb_cb_linreg.csv', index=False, float_format='%.4f')
```

[0.83597733 0.16406567]

Результат: 0.940 на public lb

1.3. Метамодель - XGBRegressor

In [32]:

```
1 %%time
2 xg = XGBRegressor()
3 xg.fit(X_train, y_train)
4 preds = xg.predict(X_val)
5 preds = np.round(np.clip(preds, 0, 10)).astype(int)
6 score = f1_score(y_val, preds, average = 'macro')
7 print(f'score on val: {score}')
```

[11:22:54] WARNING: src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.

score on val: 0.9377933690895529

Wall time: 1min 51s

In [33]:

```
1 xg.fit(oof_df, y)
2 preds = xg.predict(test_df)
3 preds = np.round(np.clip(preds, 0, 10)).astype(int)
4 sample_df = pd.read_csv("data/sample_submission.csv", dtype={'time':str})
5 sample_df['open_channels'] = preds
6 sample_df.to_csv('lgb_cb_xgb.csv', index=False, float_format='%.4f')
```

[11:25:25] WARNING: src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.

[11:25:25] WARNING: src/learner.cc:686: Tree method is automatically selected to be 'approx' for faster speed. To use old behavior (exact greedy algorithm on single machine), set tree_method to 'exact'.

Результат: 0.824 на public lb

2. Добавим oof-предсказания RandomForest

2.1. Метамодель - линейная регрессия

In [6]:

```
1 rf_oof = np.load('for_stacking/rf1_oof_preds.npy')
2 lgb_oof = np.load('for_stacking/lgb13_oof_preds.npy')
3 cb_oof = np.load('for_stacking/oof_best_catboost.npy')
```

```
In [7]: 1 oof_df = pd.DataFrame({'lgb_oof': lgb_oof, 'cb_oof': cb_oof,
2                               'rf_oof': rf_oof})
3 oof_df.head()
```

Out[7]:

	lgb_oof	cb_oof	rf_oof
0	0.132022	0.030460	0.0
1	0.097993	-0.028301	0.0
2	0.081194	-0.085979	0.0
3	0.138549	0.036849	0.0
4	0.060607	0.035619	0.0

```
In [7]: 1 train = pd.read_csv('data/train_clean.csv')
2 test = pd.read_csv('data/test_clean.csv')
3 y = train.open_channels.values
4 train = train.drop(columns=['open_channels'])
```

```
In [13]: 1 X_train, X_val, y_train, y_val = train_test_split(oof_df, y,
2                                                         random_state=17,
3                                                         test_size=0.3)
4
```

```
In [18]: 1 %%time
2 lr = LinearRegression()
3 lr.fit(X_train, y_train)
4
5 print(lr.coef_)

[0.7904124  0.16517715  0.04447739]
Wall time: 818 ms
```

```
In [19]: 1 preds = lr.predict(X_val)
2 preds = np.round(np.clip(preds, 0, 10)).astype(int)
3 score = f1_score(y_val, preds, average = 'macro')
4
5 print(f'score on val: {score}')
```

score on val: 0.9378588402608

```
In [30]: 1 scores = cv_loop(oof_df, y, LinearRegression())
2         print(scores)
3         print(np.mean(scores))
```

A Jupyter widget could not be displayed because the widget state could not be found. This could happen if the kernel storing the widget is no longer available, or if the widget state was not saved in the notebook. You may be able to create the widget by running the appropriate cells.

```
[0.9378524357477883, 0.9381910752966391, 0.9376301521633792, 0.939490083507193
2, 0.9385950090858696]
0.9383517511601738
```

3. Добавим вероятности из ноутбука по RandomForest

3.1. Метамодель - LinearRegression()

```
In [8]: 1 lgb_oof = np.load('for_stacking/lgb13_oof_preds.npy')
2         cb_oof = np.load('for_stacking/oof_best_catboost.npy')
3         rf_oof_probs = np.load('for_stacking/rf_train_probs.npy')
4
5         oof_df = pd.DataFrame({'lgb_oof': lgb_oof,
6                                'cb_oof': cb_oof})
7         for i in range(11):
8             oof_df[f'rf_oof_prob{str(i)}'] = rf_oof_probs[:, i]
```

```
In [9]: 1 scores = cv_loop(oof_df, y, LinearRegression())
2         print(scores)
3         print(np.mean(scores))
```

A Jupyter widget could not be displayed because the widget state could not be found. This could happen if the kernel storing the widget is no longer available, or if the widget state was not saved in the notebook. You may be able to create the widget by running the appropriate cells.

```
0.9379607836734138
0.9381939577977526
0.9377431802136612
0.9395857113729804
0.9387336790542435
```

```
[0.9379607836734138, 0.9381939577977526, 0.9377431802136612, 0.939585711372980
4, 0.9387336790542435]
0.9384434624224103
```

```
In [10]: 1 %%time
2 X_train, X_val, y_train, y_val = train_test_split(oof_df, y,
3                                                    random_state=17,
4                                                    test_size=0.3)
5
6 lr = LinearRegression()
7 lr.fit(X_train, y_train)
8
9 print(lr.coef_)
10
11 preds = lr.predict(X_val)
12 preds = np.round(np.clip(preds, 0, 10)).astype(int)
13 score = f1_score(y_val, preds, average = 'macro')
14
15 print(f'score on val: {score}')
```

```
[ 7.25332303e-01  1.31141734e-01 -7.20452826e-01 -5.76616182e-01
-4.33604634e-01 -2.89085199e-01 -1.45589179e-01 -3.09331149e-04
 1.39950744e-01  2.85790302e-01  4.30085623e-01  5.75450079e-01
 7.34380604e-01]
score on val: 0.9378757522818262
Wall time: 3.43 s
```

```
In [75]: 1 lgb_preds = np.load('for_stacking/lgb13_test_preds.npy')
2 cb_preds = np.load('for_stacking/preds_best_catboost.npy')
3 rf_probs = np.load('for_stacking/rf_test_probs.npy')
4
5 test_df = pd.DataFrame({'lgb_oof': lgb_preds, 'cb_oof': cb_preds})
6 for i in range(11):
7     test_df[f'rf_oof_prob{str(i)}'] = rf_probs[:, i]
```

```
In [12]: 1 lr = LinearRegression()
2 lr.fit(oof_df, y)
3
4 print(lr.coef_)
5
6 preds = lr.predict(test_df)
7 preds = np.round(np.clip(preds, 0, 10)).astype(int)
8
9 sample_df = pd.read_csv("data/sample_submission.csv", dtype={'time':str})
10 sample_df['open_channels'] = preds
11 sample_df.to_csv('stacking3.1.1.csv', index=False, float_format='%.4f')
```

```
[ 7.23415356e-01  1.32802225e-01 -1.19220250e+10 -1.19220250e+10
-1.19220250e+10 -1.19220250e+10 -1.19220250e+10 -1.19220250e+10
-1.19220250e+10 -1.19220250e+10 -1.19220250e+10 -1.19220250e+10
-1.19220250e+10]
```

3.2. Метамодел ь - XGBClassifier()

```
In [7]: 1 scores = cv_loop(oof_df, y, XGBClassifier(verbosity=1,
2                                           n_estimators=1))
3 print(scores)
4 print(np.mean(scores))
```

...

3.3. Метамодел ь - CatBoostClassifier

```
In [71]: 1 from catboost import CatBoostClassifier
```

```
In [72]: 1 cb = CatBoostClassifier(iterations=1500,
2                               verbose=1)
3 cb.fit(oof_df, y)
```

Learning rate set to 0.086731


0:	learn: 1.7197795	total: 5.87s	remaining: 2h 26m 39s
1:	learn: 1.4436677	total: 10.1s	remaining: 2h 6m 15s
2:	learn: 1.2474957	total: 14.2s	remaining: 1h 58m 20s
3:	learn: 1.0976631	total: 18.4s	remaining: 1h 54m 32s
4:	learn: 0.9609647	total: 22.5s	remaining: 1h 52m 10s
5:	learn: 0.8643736	total: 26.6s	remaining: 1h 50m 25s
6:	learn: 0.7827403	total: 30.8s	remaining: 1h 49m 34s
7:	learn: 0.7139835	total: 35.1s	remaining: 1h 49m
8:	learn: 0.6556434	total: 39.4s	remaining: 1h 48m 42s
9:	learn: 0.6031464	total: 43.7s	remaining: 1h 48m 27s
10:	learn: 0.5563732	total: 47.9s	remaining: 1h 48m 5s
11:	learn: 0.5158112	total: 52.2s	remaining: 1h 47m 56s
12:	learn: 0.4792302	total: 56.4s	remaining: 1h 47m 31s
13:	learn: 0.4472814	total: 1m	remaining: 1h 47m 26s
14:	learn: 0.4181340	total: 1m 4s	remaining: 1h 47m 12s
15:	learn: 0.3918102	total: 1m 9s	remaining: 1h 47m 8s
16:	learn: 0.3652644	total: 1m 13s	remaining: 1h 47m 8s
17:	learn: 0.3438955	total: 1m 18s	remaining: 1h 47m 6s
18:	learn: 0.3245000	total: 1m 22s	remaining: 1h 47m

```
In [79]: 1 preds = cb.predict(test_df)
2 preds = pred_proc(preds)
3 sample_df = pd.read_csv("data/sample_submission.csv", dtype={'time':str})
4 sample_df['open_channels'] = preds
5 sample_df.to_csv("stacking3.3.csv", index=False, float_format='%.4f')
```

4. Разный блендинг на разных батчах



```
In [31]: 1 plt.plot(np.arange(0, len(train.signal.values)), train.signal.values,  
2              alpha=0.5)
```

Out[31]: [<matplotlib.lines.Line2D at 0x1de67580608>]



```
In [32]: 1 plt.plot(np.arange(0, len(test.signal.values)), test.signal.values,  
2              alpha=0.5)
```

Out[32]: [<matplotlib.lines.Line2D at 0x1de675e2c48>]



```
In [26]: 1 for i in range(10):
2         print(np.mean(train.signal.values[i*500000:(i+1)*500000]))
```

-2.681406359199999
-2.6831050761904818
-1.8131248770000001
-0.09298384679999996
3.3578080268
1.6808430502000005
-1.8039282956274345
-0.10786146542743458
1.658963467372566
3.3109071119725657

```
In [30]: 1 for i in range(4):
2         print(np.mean(test.signal.values[i*500000:(i+1)*500000]))
```

-1.076758010171455
1.119842871828547
-2.5992788104274345
-2.6210051965999996

```
In [28]: 1 preds = np.zeros(2*10**6)
```

Для первого батча из теста подбираем коэффициенты по 3, 4, 7, 8 батчам из трейна

```
In [47]: 1 train34 = oof_df.iloc[np.arange(2*500000, 4*500000), ]
2 train78 = oof_df.iloc[np.arange(6*500000, 8*500000), ]
3 train3478 = pd.concat([train34, train78])
4 y3478 = y[list(np.arange(2*500000, 4*500000)) + list(np.arange(6*500000, 8*500000))]
```

```
In [54]: 1 lr = LinearRegression()
2 lr.fit(train3478, y3478)
3 print(lr.coef_)
4 preds3478 = lr.predict(test_df.iloc[np.arange(0, 500000)])
5 preds[np.arange(0, 500000)] = preds3478
```

```
[ 6.52878545e-01  9.44760369e-02 -1.73687652e+10 -1.73687652e+10
 -1.73687652e+10 -1.73687652e+10 -1.73687652e+10 -1.73687652e+10
 -1.73687652e+10 -1.73687652e+10 -1.73687652e+10 -1.73687652e+10
 -1.73687652e+10]
```

Для второго батча из теста подбираем коэффициенты по 5,6,9,10 батчам из трейна

```
In [56]: 1 train56 = oof_df.iloc[np.arange(4*500000, 6*500000), ]
2 train910 = oof_df.iloc[np.arange(8*500000, 10*500000), ]
3 train5691 = pd.concat([train56, train910])
4 y5691 = y[list(np.arange(4*500000, 6*500000)) + list(np.arange(8*500000, 10*
5 lr = LinearRegression()
6 lr.fit(train5691, y5691)
7 print(lr.coef_)
8 preds5691 = lr.predict(test_df.iloc[np.arange(500000, 2*500000)])
9 preds[np.arange(500000, 2*500000)] = preds5691
```

```
[ 0.73059474  0.14162184 -0.67248281 -0.51879788 -0.38423512 -0.2542972
-0.12451996  0.00411493  0.12924003  0.25904027  0.38739255  0.51733709
 0.6572081 ]
```

Для третьего и четвертого батча из теста по 1 и 2 батчам из трейна

```
In [57]: 1 train12 = oof_df.iloc[np.arange(0*500000, 2*500000), ]
2 y12 = y[list(np.arange(0*500000, 2*500000))]
3 lr = LinearRegression()
4 lr.fit(train12, y12)
5 print(lr.coef_)
6 preds12 = lr.predict(test_df.iloc[np.arange(2*500000, 4*500000)])
7 preds[np.arange(2*500000, 4*500000)] = preds12
```

```
[ 3.93108197e-01  7.63297246e-02 -1.42249995e+10 -1.42249995e+10
-1.42249995e+10 -1.42249995e+10 -1.42249996e+10 -1.42249995e+10
-1.42249900e+10  0.00000000e+00  0.00000000e+00  0.00000000e+00
 0.00000000e+00]
```

```
In [59]: 1 preds = pred_proc(preds)
2 sample_df = pd.read_csv("data/sample_submission.csv", dtype={'time':str})
3 sample_df['open_channels'] = preds
4 sample_df.to_csv("stacking4.csv", index=False, float_format='%.4f')
```

Результат: 0.938 на lb

5. Добавим результаты алгоритма Витерби

5.1. Мета модель - линейная регрессия

```

In [78]: 1 # исходные данные
2 train = pd.read_csv('data/train_clean.csv')
3 test = pd.read_csv('data/test_clean.csv')
4 y = train.open_channels.values
5 train = train.drop(columns=['open_channels'])
6
7
8 # трейн
9 lgb_oof = np.load('for_stacking/lgb13_oof_preds.npy')
10 cb_oof = np.load('for_stacking/oof_best_catboost.npy')
11 rf_oof_probs = np.load('for_stacking/rf_train_probs.npy')
12 vit_oof_preds = np.load('for_stacking/oof_viterbi.npy')
13 oof_df = pd.DataFrame({'lgb_oof': lgb_oof,
14                        'cb_oof': cb_oof,
15                        'viterbi': vit_oof_preds})
16 for i in range(11):
17     oof_df[f'rf_oof_prob{str(i)}'] = rf_oof_probs[:, i]
18
19
20 # тест
21 lgb_preds = np.load('for_stacking/lgb13_test_preds.npy')
22 cb_preds = np.load('for_stacking/preds_best_catboost.npy')
23 rf_probs = np.load('for_stacking/rf_test_probs.npy')
24 vit_preds = np.load('for_stacking/preds_viterbi.npy')
25 test_df = pd.DataFrame({'lgb_oof': lgb_preds,
26                        'cb_oof': cb_preds,
27                        'viterbi': vit_preds})
28 for i in range(11):
29     test_df[f'rf_oof_prob{str(i)}'] = rf_probs[:, i]

```

```

In [68]: 1 scores = cv_loop(oof_df, y, LinearRegression())
2 print(scores)
3 print(np.mean(scores))

```

A Jupyter widget could not be displayed because the widget state could not be found. This could happen if the kernel storing the widget is no longer available, or if the widget state was not saved in the notebook. You may be able to create the widget by running the appropriate cells.

```

0.9379437431519608
0.9381737172630964
0.9377425770015045
0.9395930654619
0.9387499979561834

```

```

[0.9379437431519608, 0.9381737172630964, 0.9377425770015045, 0.9395930654619,
0.9387499979561834]
0.938440620166929

```

```

In [69]: 1 lr = LinearRegression()
2 lr.fit(oof_df, y)

```

Out[69]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)

In [70]:

1	lr.coef_
---	----------

Out[70]: array([7.23245139e-01, 1.32816043e-01, -4.42761636e-04, -7.24394528e-01,
-5.79622293e-01, -4.35792810e-01, -2.90513810e-01, -1.46104082e-01,
-5.90173974e-04, 1.40557225e-01, 2.87431766e-01, 4.32618197e-01,
 5.79539631e-01, 7.36870877e-01])