

В этом ноутбуке происходят попытки выбить максимальный скор из случайного леса. Discussion на kaggle показывают, что лес здесь неплохо справляется

```
In [1]: 1 import numpy as np
2 import matplotlib.pyplot as plt
3 import pandas as pd
4 from tqdm import tqdm_notebook
5 from sklearn.ensemble import RandomForestClassifier
6 from sklearn.base import BaseEstimator, TransformerMixin
7 from sklearn.model_selection import train_test_split, KFold
8 from sklearn.metrics import f1_score
9 import gc
```

0. Скрипты для загрузки данных, обучения, кросс-валидации, загрузки предсказаний в файл

Загрузка данных

```
In [2]: 1 from feature_engineering import reduce_mem_usage, add_rolling_features
2 from feature_engineering import exponential_smoothing, signal_shifts
3 from feature_engineering import batch_stats2, add_minus_signal
4 from feature_engineering import delete_objects_after_rolling
5 from feature_engineering import add_quantiles, add_target_encoding
```

```
In [3]: 1 def prepare_df(df, shifts):
2     df = reduce_mem_usage(df)
3     df = signal_shifts(df, shifts)
4     df = reduce_mem_usage(df)
5
6     if 'open_channels' in df.columns:
7         y = df['open_channels']
8         df = df.drop(columns=['time'])
9         return df, y
10    else:
11        df = df.drop(columns=['time'])
12        return df
```

```

In [4]: 1 train = pd.read_csv('data/train_clean.csv')
        2 test = pd.read_csv('data/test_clean.csv')
        3
        4 shifts = list(np.arange(-20, 0)) + list(np.arange(1, 21))
        5
        6 train["category"] = 0
        7 test["category"] = 0
        8
        9 # train segments with more then 9 open channels classes
       10 train.loc[2_000_000:2_500_000-1, 'category'] = 1
       11 train.loc[4_500_000:5_000_000-1, 'category'] = 1
       12
       13 # test segments with more then 9 open channels classes (potentially)
       14 test.loc[500_000:600_000-1, "category"] = 1
       15 test.loc[700_000:800_000-1, "category"] = 1
       16
       17 X_train, y_train = prepare_df(train, shifts)
       18 X_test = prepare_df(test, shifts)
       19
       20 y_train = np.array(y_train)
       21
       22 # add_quantiles(X_train, X_test, [3, 7, 15])
       23 # add_target_encoding(X_train, X_test, [3, 7, 15])
       24
       25 X_train = reduce_mem_usage(X_train)
       26 X_test = reduce_mem_usage(X_test)
       27
       28 X_train = X_train.drop(columns=['open_channels'])

```

Mem. usage decreased to 28.61 Mb (81.2% reduction)
 Mem. usage decreased to 410.08 Mb (14.0% reduction)
 Mem. usage decreased to 9.54 Mb (79.2% reduction)
 Mem. usage decreased to 162.12 Mb (7.6% reduction)
 Mem. usage decreased to 400.54 Mb (0.0% reduction)
 Mem. usage decreased to 158.31 Mb (0.0% reduction)

Обучение леса

```

In [5]: 1 def fit_model(X_train, y_train, params):
2         gc.collect()
3
4         print('splitting...')
5         X_train_, X_valid, y_train_, y_valid = train_test_split(X_train, y_train,
6                                                                 test_size=0.3,
7                                                                 random_state=17)
8
9         print('fit...')
10        model = RandomForestClassifier(**params)
11        model.fit(X_train_, y_train_)
12
13        print('predict...')
14        prediction = model.predict(X_valid)
15        score = f1_score(y_valid, prediction, average = 'macro')
16
17        print(f'score = {score}')
18
19        return model
20
21 def fit_model_with_save(X_train, y_train, X_test, params, modelname):
22     gc.collect()
23
24     print('fit...')
25     model = RandomForestClassifier(**params)
26     model.fit(X_train, y_train)
27
28     print('predict...')
29     prediction = model.predict(X_test)
30     np.save(modelname + '_test_preds.npy', prediction)
31
32     print('saving predictions...')
33     sample_df = pd.read_csv("data/sample_submission.csv", dtype={'time':str})
34     sample_df['open_channels'] = prediction
35     sample_df.to_csv(modelname + '.csv', index=False, float_format='%.4f')
36
37     print('probability...')
38     probs = model.predict_proba(X_test)
39     np.save(modelname + '_test_probs.npy', probs)
40
41     return model

```

Кросс-валидация

```

In [6]: 1 def cv_loop(X_train, y_train, X_test, params, modelname):
2         n_fold = 5
3         folds = KFold(n_splits=n_fold, shuffle=True, random_state=17)
4
5         oof = np.zeros(X_train.shape[0])
6         oof_probs = np.zeros((X_train.shape[0], 11))
7
8         prediction = np.zeros(X_test.shape[0])
9         scores = []
10
11     for training_index, validation_index in tqdm_notebook(folds.split(X_train)):
12         gc.collect()
13
14         # разбиение на трэйн и валидацию
15         X_train_ = X_train.iloc[training_index]
16         y_train_ = y_train[training_index]
17         X_valid = X_train.iloc[validation_index]
18         y_valid = y_train[validation_index]
19
20         # обучение модели
21         model = RandomForestClassifier(**params)
22         model.fit(X_train_, y_train_)
23
24         # скор на валидации
25         preds = model.predict(X_valid)
26         oof[validation_index] = preds
27         score = f1_score(y_valid, preds, average = 'macro')
28         scores.append(score)
29
30         # вероятности на валидации
31         probs = model.predict_proba(X_valid)
32         oof_probs[validation_index] = probs
33
34         # предсказание на тесте
35         preds = model.predict(X_test)
36         prediction += preds
37
38         print(f'score: {score}')
39
40     prediction /= n_fold
41     prediction = np.round(np.clip(prediction, 0, 10)).astype(int)
42
43     np.save(modelname + '_cv_test_preds.npy', prediction)
44     np.save(modelname + '_oof_preds.npy', oof)
45     np.save(modelname + '_oof_probs.npy', oof_probs)
46
47     sample_df = pd.read_csv("data/sample_submission.csv", dtype={'time':str})
48     sample_df['open_channels'] = prediction
49     sample_df.to_csv(modelname + '_cv.csv', index=False, float_format='%.4f')
50
51     return scores, oof, prediction

```

1. Бейзлайн

Параметры для RF и для признаков возьмем из этого ноутбука:

<https://www.kaggle.com/sggpls/shifted-rfc-pipeline> (<https://www.kaggle.com/sggpls/shifted-rfc-pipeline>)

```
In [7]: 1 %%time
2
3 params = {
4     'n_estimators': 200,
5     'max_depth': 19,
6     'max_features': 10,
7     'random_state': 17,
8     'n_jobs': -1,
9     'verbose': 2
10 }
11
12 scores, oof, prediction = cv_loop(X_train, y_train, X_test, params, 'rf1')
```

A Jupyter widget could not be displayed because the widget state could not be found. This could happen if the kernel storing the widget is no longer available, or if the widget state was not saved in the notebook. You may be able to create the widget by running the appropriate cells.

[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 8 concurrent workers.

```
In [8]: 1 print(np.mean(scores))
```

0.9377745874531744

In [7]:

```
1 %%time
2 params = {
3     'n_estimators': 200,
4     'max_depth': 19,
5     'max_features': 10,
6     'random_state': 17,
7     'n_jobs': -1,
8     'verbose': 2
9 }
10 forest = fit_model_with_save(X_train, y_train, X_test,
11                               params, 'rf1')
```

fit...

[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 8 concurrent workers.

building tree 1 of 200building tree 2 of 200building tree 3 of 200building tree 4 of 200building tree 5 of 200building tree 6 of 200building tree 7 of 200

building tree 8 of 200

building tree 9 of 200building tree 10 of 200

building tree 11 of 200
building tree 12 of 200
building tree 13 of 200
building tree 14 of 200

Результат: 0.939 на public lb.

2. Добавим таргет энкодинг, подкрутим параметры

In [11]:

```
1 %%time
2
3 params = {
4     'n_estimators': 500,
5     'max_depth': 25,
6     'max_features': 15,
7     'random_state': 17,
8     'n_jobs': -1,
9     'verbose': 2
10 }
11
12 forest = fit_model(X_train, y_train, params)
```

splitting...

fit...

[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 8 concurrent workers.

building tree 1 of 500building tree 2 of 500building tree 3 of 500building tree 4 of 500building tree 5 of 500building tree 6 of 500building tree 7 of 500building tree 8 of 500

building tree 9 of 500
building tree 10 of 500
building tree 11 of 500
building tree 12 of 500

In [12]:

```
1 for feature, imp in zip(X_train.columns, forest.feature_importances_):  
2     print(feature, imp)
```

```
signal 0.2810031591309477  
category 0.10557577662222818  
shift_-20 0.004632864496387041  
shift_-19 0.003406895845358958  
shift_-18 0.0027048016027770012  
shift_-17 0.0020497408946804035  
shift_-16 0.001785484209258236  
shift_-15 0.0016448296427000225  
shift_-14 0.0014370097477635424  
shift_-13 0.0013014553553040565  
shift_-12 0.0013215275808369784  
shift_-11 0.0012317645641277744  
shift_-10 0.0011798043295531532  
shift_-9 0.0012400490906433485  
shift_-8 0.0012042642453490972  
shift_-7 0.0013310274181537538  
shift_-6 0.0015851816879912362  
shift_-5 0.001878256879192865  
shift_-4 0.0024917497499253078  
shift_-3 0.004939598532460965  
shift_-2 0.01272673147324323  
shift_-1 0.04282155384271084  
shift_1 0.030177157010225886  
shift_2 0.007917132601380212  
shift_3 0.004828813324772249  
shift_4 0.0022441743662006013  
shift_5 0.0020569064093582576  
shift_6 0.0017555993633171147  
shift_7 0.0019812042905531916  
shift_8 0.001295086443712815  
shift_9 0.0012752507007509001  
shift_10 0.0012460415743923575  
shift_11 0.001308760799017611  
shift_12 0.001293529238323259  
shift_13 0.0013616279430836787  
shift_14 0.0014013848078526256  
shift_15 0.0015343370034235388  
shift_16 0.0017743871197682972  
shift_17 0.0018204118122319211  
shift_18 0.002639553591744839  
shift_19 0.0033765028438492535  
shift_20 0.004378495921514019  
quant_3 0.0016559186504529312  
quant_7 0.050809769163690184  
quant_15 0.10270333776499588  
quant_3_mean 0.0007967650537465773  
quant_3_std 0.002782910556776887  
quant_3_var 0.0015300683173072784  
quant_7_mean 0.05826341553032381  
quant_7_std 0.04388602205179024  
quant_7_var 0.04973625940031991  
quant_15_mean 0.10956949554724996
```


quant_15_std 0.01152907906164337
quant_15_var 0.011577074794636813

